

LLM Judge Analysis Report: Gemini Pro 3 Preview on BBH

Generated: 2025-12-27 14:18:44

Executive Summary

Metric	Value
Total Tasks Analyzed	27
Total Samples	6,510
Overall Accuracy	60.1% (3,911/6,510)
Best Performing Task	Boolean Expressions (85.6%)
Worst Performing Task	Dyck Languages (16.1%)

Key Findings

1. **wrong_logic** is the dominant error type at 38.8% of all samples
 2. **Strong correlation between accuracy and reasoning scores** - high accuracy tasks average ≥ 4.0 reasoning score
 3. **Task difficulty varies significantly** - 16.1% to 85.6% accuracy range
- [!IMPORTANT] The model demonstrates consistent reasoning structure but makes logical/computational errors in ~39% of cases.
-

Visualizations

Task Accuracy Chart

Error Distribution

Reasoning Score Distribution

Accuracy vs Reasoning Score

Error Type Heatmap

Task Performance Details

All Tasks (Sorted by Accuracy)

Rank	Task	Total	Correct	Accuracy	Avg Score
1	Boolean Expressions	250	214	[HIGH] 85.6%	4.57
2	Movie Recommendation	250	202	[HIGH] 80.8%	4.44
3	Penguin In A Table	146	110	[HIGH] 75.3%	4.13
4	Object Counting	250	188	[HIGH] 75.2%	4.28
5	Navigate	250	185	[HIGH] 74.0%	4.18

Rank	Task	Total	Correct	Accuracy	Avg Score
6	Date Understanding	250	180	[HIGH] 72.0%	4.00
7	Sports Understanding	250	178	[HIGH] 71.2%	4.21
8	Reasoning About Colored Objects	250	175	[HIGH] 70.0%	4.16
9	Logical Deduction Three Objact	250	174	[MED] 69.6%	4.19
10	Logical Deduction Seven Objects	250	163	[MED] 65.2%	3.96
11	Web Of Lies	250	163	[MED] 65.2%	3.72
12	Word Sorting	250	162	[MED] 64.8%	3.72
13	Snarks	178	112	[MED] 62.9%	3.72
14	Formal Fallacies	250	155	[MED] 62.0%	3.48
15	Multistep Arithmetic Two	250	154	[MED] 61.6%	4.42
16	Hyperbaton	250	144	[MED] 57.6%	3.31
17	Logical Deduction Five Objects	250	143	[MED] 57.2%	3.71
18	Ruin Names	250	143	[MED] 57.2%	3.32
19	Tracking Shuffled Objects Three Objects	250	143	[MED] 57.2%	3.60
20	Causal Judgement	187	104	[MED] 55.6%	3.62
21	Tracking Shuffled Objects Five Objects	250	136	[MED] 54.4%	3.18
22	Tracking Shuffled Objects Seven Objects	250	134	[MED] 53.6%	3.10
23	Temporal Sequence	250	124	[LOW] 49.6%	2.98
24	Geometric Shapes	250	103	[LOW] 41.2%	2.65
25	Disambiguation Qa	250	99	[LOW] 39.6%	3.22
26	Salient Translation Error Detection	250	83	[LOW] 33.2%	3.03
27	Dyck Languages	249	40	[LOW] 16.1%	2.41

Error Analysis

Global Error Distribution

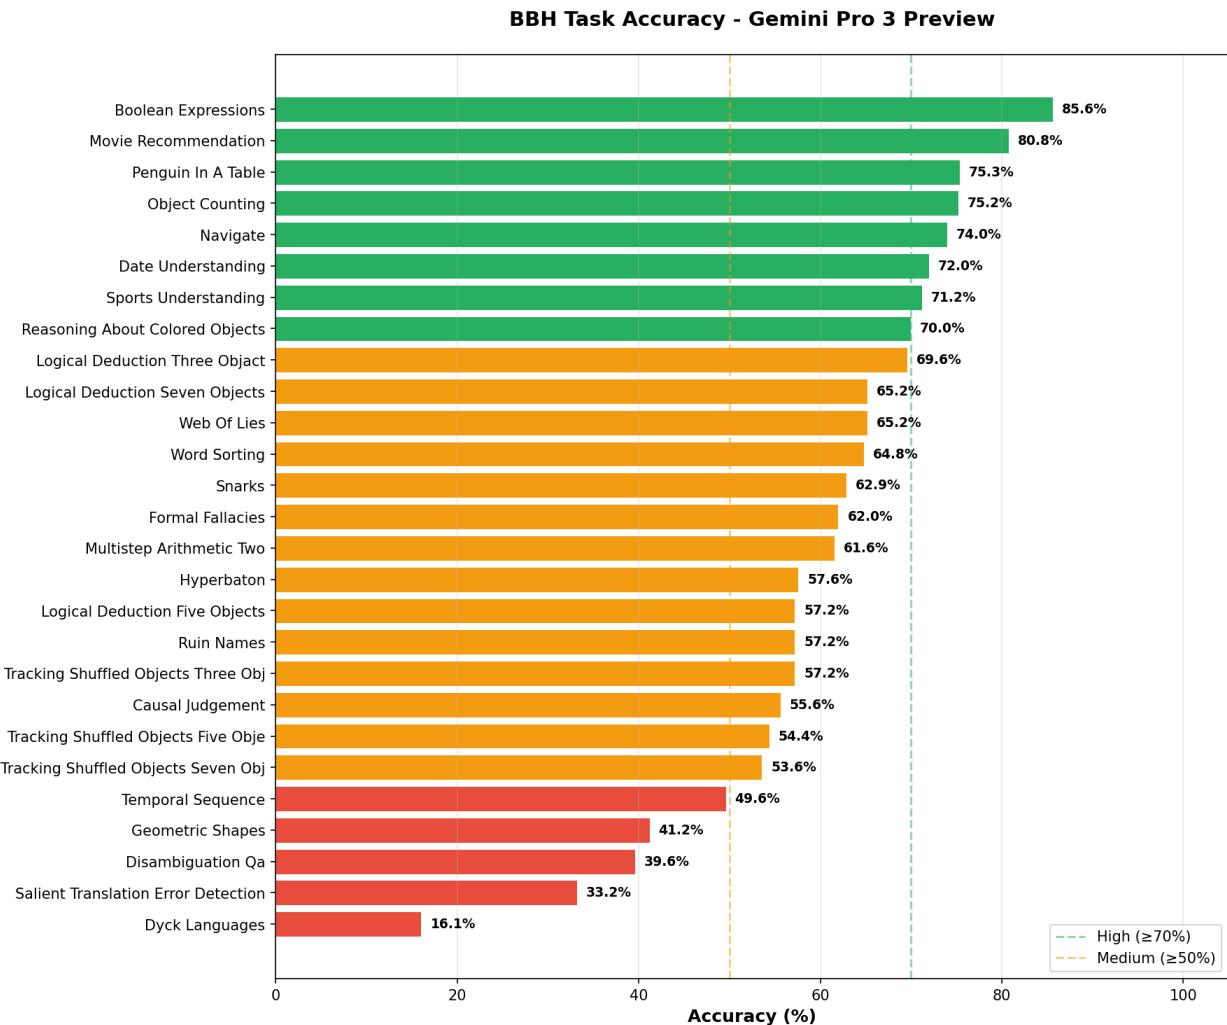


Figure 1: Task Accuracy Chart

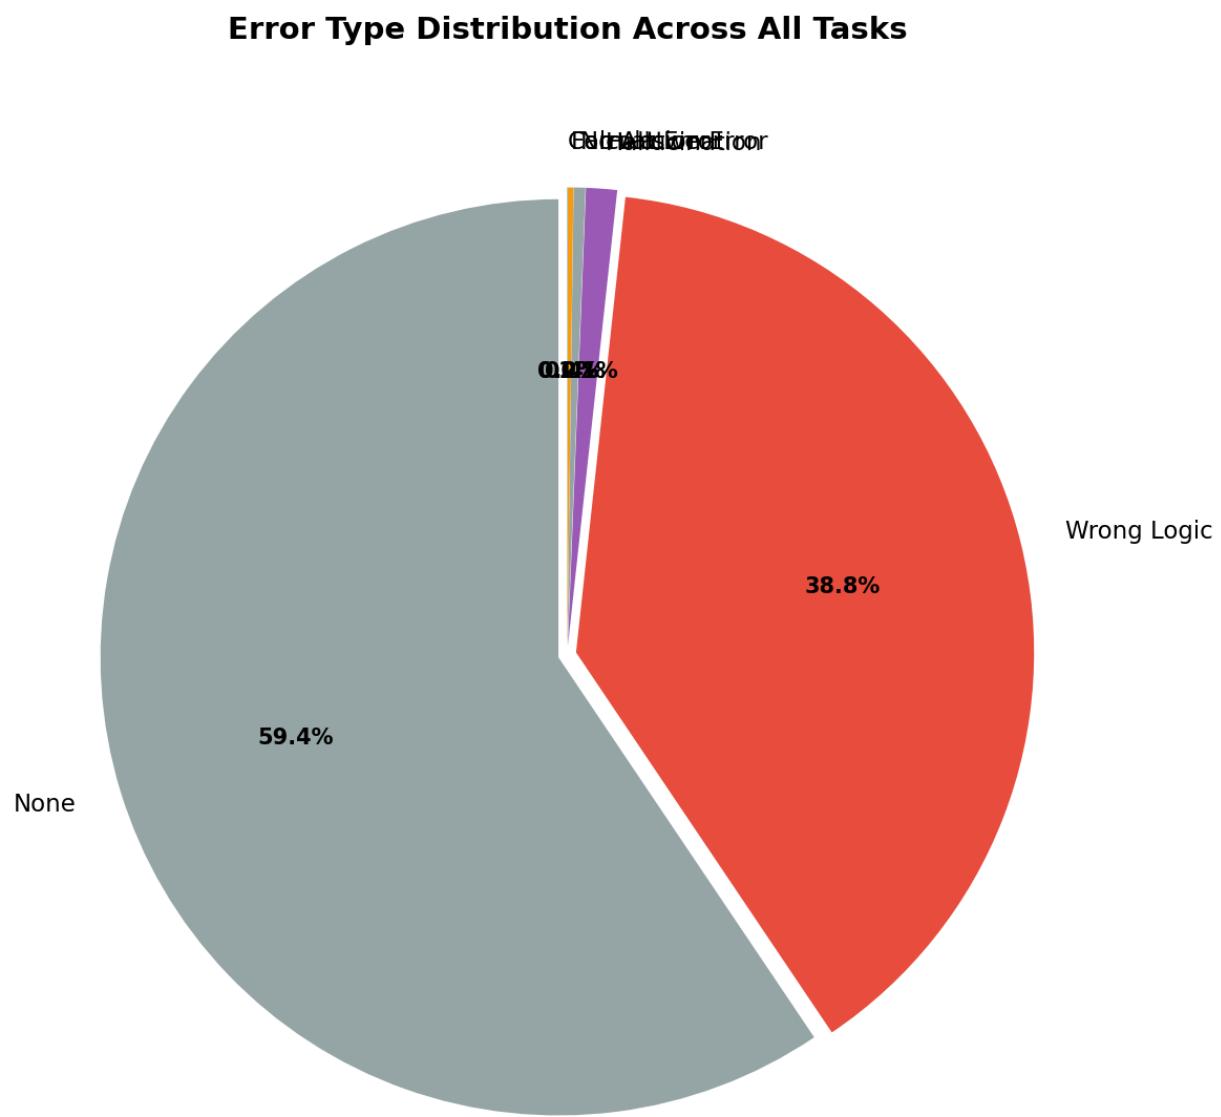


Figure 2: Error Distribution

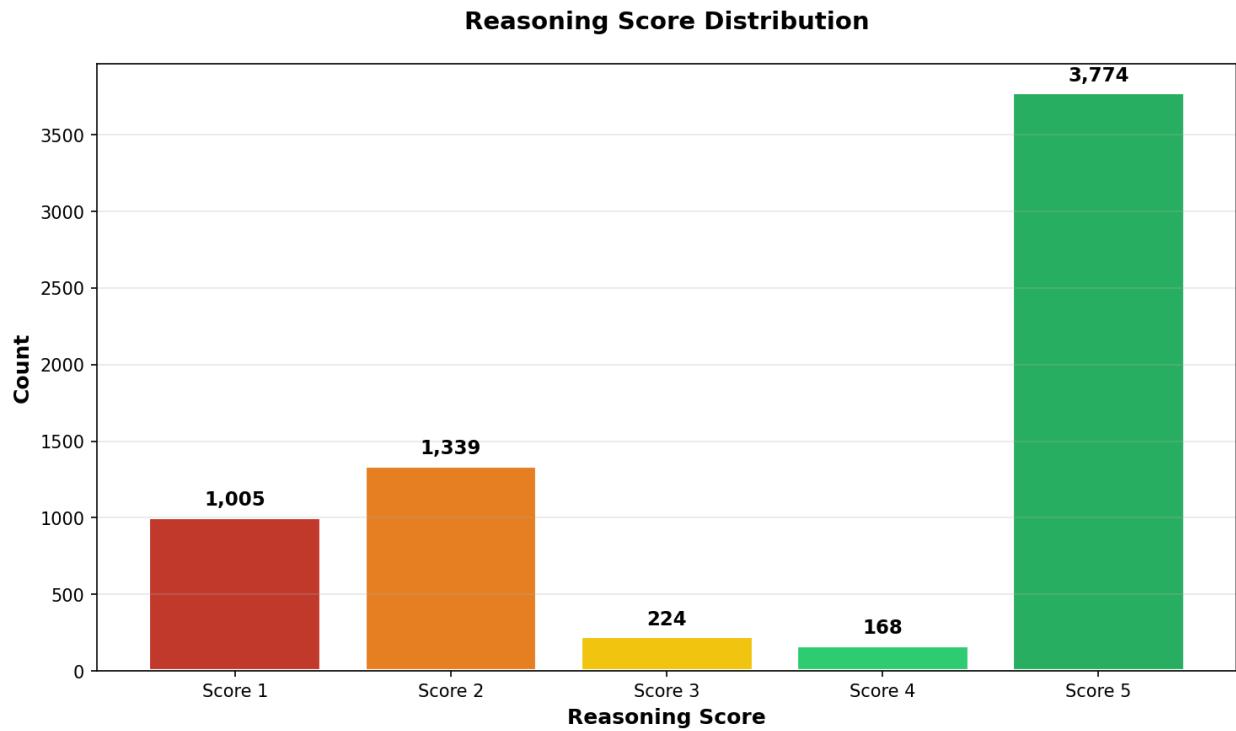


Figure 3: Reasoning Score Distribution

Error Type	Count	Percentage
none	3,869	59.4%
wrong_logic	2,529	38.8%
hallucination	71	1.1%
no_answer	27	0.4%
format_error	13	0.2%
calculation_error	1	0.0%

Reasoning Score Distribution

Score	Count	Percentage	Interpretation
5	3,774	58.0%	Perfect logic
4	168	2.6%	Minor issues
3	224	3.4%	Moderate errors
2	1,339	20.6%	Significant errors
1	1,005	15.4%	Fundamental failure

Per-Task Error Breakdown

Boolean Expressions

- **Accuracy:** 85.6% (214/250)
- **Avg Reasoning Score:** 4.57
- **Error Types:** none: 214, wrong_logic: 36

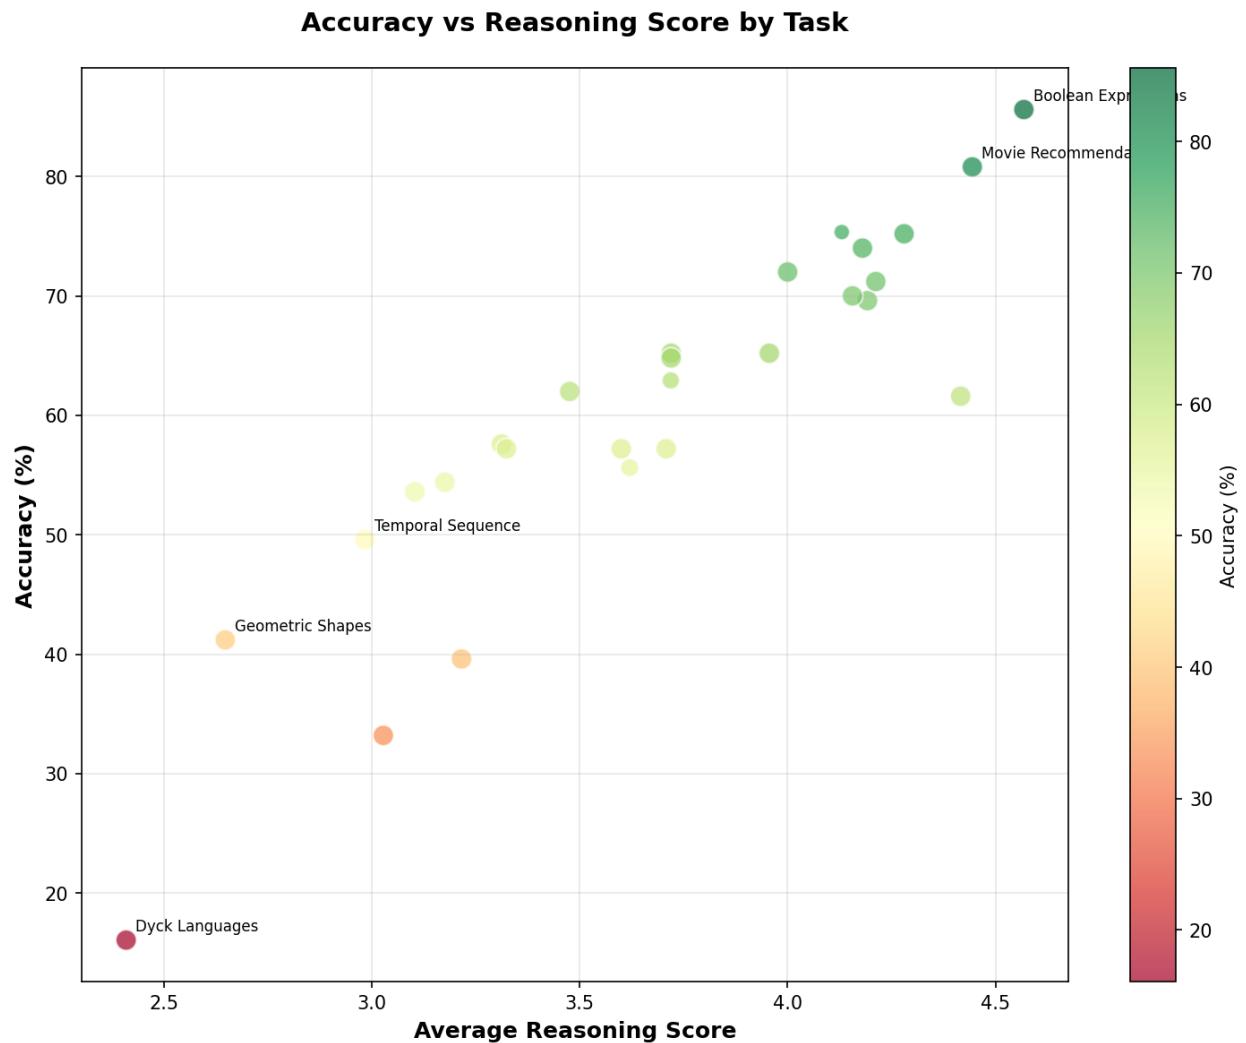


Figure 4: Accuracy vs Reasoning Score

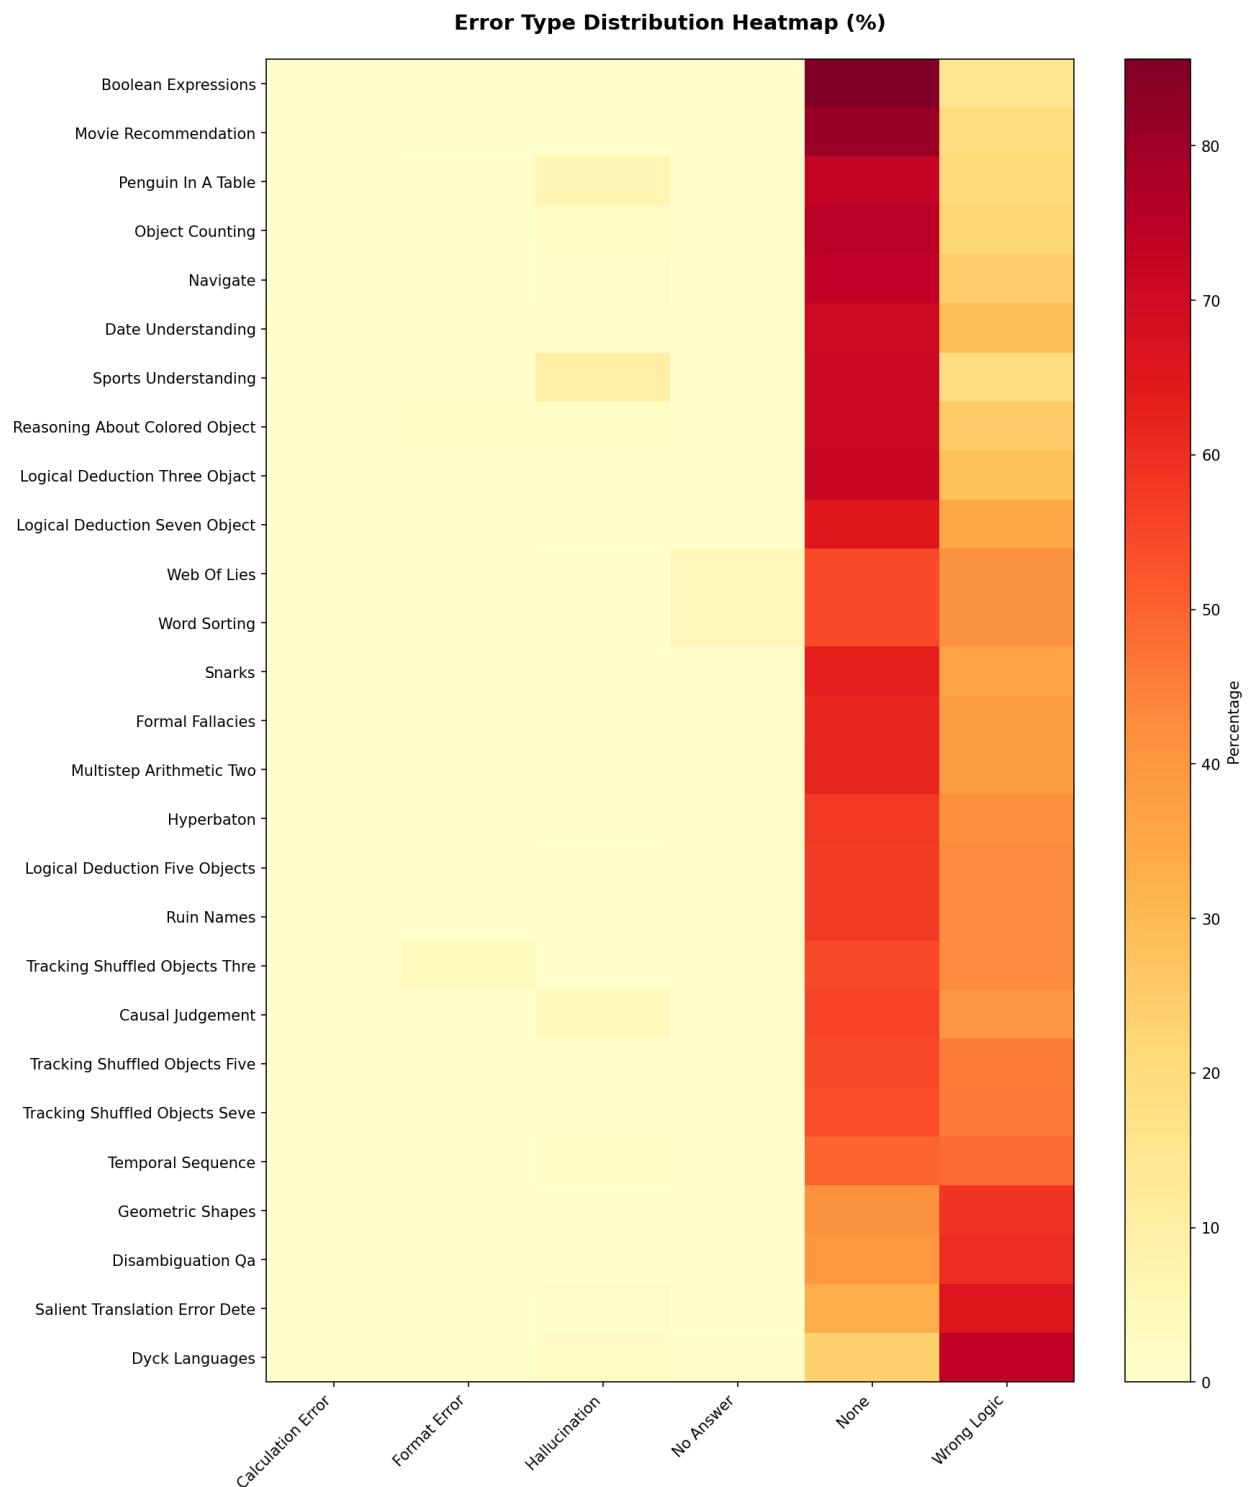


Figure 5: Error Type Heatmap

Movie Recommendation

- **Accuracy:** 80.8% (202/250)
- **Avg Reasoning Score:** 4.44
- **Error Types:** none: 202, wrong_logic: 48

Penguin In A Table

- **Accuracy:** 75.3% (110/146)
- **Avg Reasoning Score:** 4.13
- **Error Types:** none: 106, wrong_logic: 30, hallucination: 9, format_error: 1

Object Counting

- **Accuracy:** 75.2% (188/250)
- **Avg Reasoning Score:** 4.28
- **Error Types:** none: 188, wrong_logic: 56, hallucination: 6

Navigate

- **Accuracy:** 74.0% (185/250)
- **Avg Reasoning Score:** 4.18
- **Error Types:** none: 185, wrong_logic: 62, hallucination: 3

Date Understanding

- **Accuracy:** 72.0% (180/250)
- **Avg Reasoning Score:** 4.00
- **Error Types:** none: 176, wrong_logic: 72, calculation_error: 1, no_answer: 1

Sports Understanding

- **Accuracy:** 71.2% (178/250)
- **Avg Reasoning Score:** 4.21
- **Error Types:** none: 178, wrong_logic: 47, hallucination: 25

Reasoning About Colored Objects

- **Accuracy:** 70.0% (175/250)
- **Avg Reasoning Score:** 4.16
- **Error Types:** none: 178, wrong_logic: 63, format_error: 5, hallucination: 4

Logical Deduction Three Object

- **Accuracy:** 69.6% (174/250)
- **Avg Reasoning Score:** 4.19
- **Error Types:** none: 179, wrong_logic: 71

Logical Deduction Seven Objects

- **Accuracy:** 65.2% (163/250)
- **Avg Reasoning Score:** 3.96
- **Error Types:** none: 163, wrong_logic: 87

Web Of Lies

- **Accuracy:** 65.2% (163/250)
- **Avg Reasoning Score:** 3.72
- **Error Types:** none: 136, wrong_logic: 102, no_answer: 11, hallucination: 1

Word Sorting

- **Accuracy:** 64.8% (162/250)
- **Avg Reasoning Score:** 3.72
- **Error Types:** none: 136, wrong_logic: 102, no_answer: 11, hallucination: 1

Snarks

- **Accuracy:** 62.9% (112/178)
- **Avg Reasoning Score:** 3.72
- **Error Types:** none: 112, wrong_logic: 65, no_answer: 1

Formal Fallacies

- **Accuracy:** 62.0% (155/250)
- **Avg Reasoning Score:** 3.48
- **Error Types:** none: 155, wrong_logic: 95

Multistep Arithmetic Two

- **Accuracy:** 61.6% (154/250)
- **Avg Reasoning Score:** 4.42
- **Error Types:** none: 154, wrong_logic: 95, no_answer: 1

Hyperbaton

- **Accuracy:** 57.6% (144/250)
- **Avg Reasoning Score:** 3.31
- **Error Types:** none: 144, wrong_logic: 105, hallucination: 1

Logical Deduction Five Objects

- **Accuracy:** 57.2% (143/250)
- **Avg Reasoning Score:** 3.71
- **Error Types:** none: 143, wrong_logic: 107

Ruin Names

- **Accuracy:** 57.2% (143/250)
- **Avg Reasoning Score:** 3.32
- **Error Types:** none: 143, wrong_logic: 107

Tracking Shuffled Objects Three Objects

- **Accuracy:** 57.2% (143/250)
- **Avg Reasoning Score:** 3.60
- **Error Types:** none: 136, wrong_logic: 107, format_error: 7

Causal Judgement

- **Accuracy:** 55.6% (104/187)
- **Avg Reasoning Score:** 3.62
- **Error Types:** none: 104, wrong_logic: 75, hallucination: 8

Tracking Shuffled Objects Five Objects

- **Accuracy:** 54.4% (136/250)
- **Avg Reasoning Score:** 3.18
- **Error Types:** none: 136, wrong_logic: 114

Tracking Shuffled Objects Seven Objects

- **Accuracy:** 53.6% (134/250)
- **Avg Reasoning Score:** 3.10
- **Error Types:** none: 134, wrong_logic: 116

Temporal Sequence

- **Accuracy:** 49.6% (124/250)
- **Avg Reasoning Score:** 2.98
- **Error Types:** none: 124, wrong_logic: 121, hallucination: 4, no_answer: 1

Geometric Shapes

- **Accuracy:** 41.2% (103/250)
- **Avg Reasoning Score:** 2.65
- **Error Types:** wrong_logic: 147, none: 103

Disambiguation Qa

- **Accuracy:** 39.6% (99/250)
- **Avg Reasoning Score:** 3.22
- **Error Types:** wrong_logic: 151, none: 99

Salient Translation Error Detection

- **Accuracy:** 33.2% (83/250)
- **Avg Reasoning Score:** 3.03
- **Error Types:** wrong_logic: 164, none: 83, hallucination: 3

Dyck Languages

- **Accuracy:** 16.1% (40/249)
 - **Avg Reasoning Score:** 2.41
 - **Error Types:** wrong_logic: 184, none: 58, hallucination: 6, no_answer: 1
-

Sample Cases

Below are randomly sampled correct and incorrect cases from select tasks.

Dyck Languages

Incorrect Cases:

- **ID 229:** Answer=[< { < (())) > } > , Score=2, Error=wrong_logic
 - Model added]]. GT is]]. Correct.
- **ID 108:** Answer=< { < ())) , Score=1, Error=wrong_logic
 - Model modified structure.

Correct Cases:

- **ID 71:** Answer={ [< [< { < { } > } > ((< { [] } > { { } }) { })] >] }, Score=5^c
- **ID 19:** Answer=(()) [{ (< >) }], Score=5^c

Salient Translation Error Detection

Incorrect Cases:

- **ID 180:** Answer=A, Score=2, Error=wrong_logic
 - The translation changed ‘kleines’ (small) to ‘big’. This is an Antonym error (C). Model selected (A).
- **ID 90:** Answer=D, Score=2, Error=wrong_logic
 - The translation changed ‘deutsche’ to ‘French’. This is a Modifier/Adjective error (A). Model selected (D).

Correct Cases:

- **ID 99:** Answer=B, Score=5^c
- **ID 153:** Answer=E, Score=5^c

Boolean Expressions

Incorrect Cases:

- **ID 178:** Answer=True, Score=2, Error=wrong_logic
 - Precedence error. True and (not not False) and (not False). True and False and True -> False. Model got True.
- **ID 222:** Answer=True, Score=2, Error=wrong_logic
 - Precedence error. False or ((not False) and False). True and False is False. False or False is False. Model said True.

Correct Cases:

- **ID 82:** Answer=True, Score=5^c
 - **ID 155:** Answer=False, Score=5^c
-

Methodology

LLM Judge Configuration

The evaluation was performed using **Gemini Pro 3 Preview** as an LLM-as-Judge with the following approach:

1. **Semantic Matching:** Prioritizes the semantic meaning of the final conclusion over strict formatting
2. **Answer Extraction:** Ignores <think> traces, focuses on final response text
3. **Equivalence Rules:** Maps various answer formats to ground truth:
 - Multiple Choice: (A) = A, Option A, [A], \boxed{A}
 - Boolean: True = true, yes, valid; False = false, no, invalid

- Plausibility: `yes` = plausible; `no` = implausible

Scoring Rubric

Score	Meaning
5	Perfect logic and correct answer
4	Minor issues in reasoning
3	Moderate errors
2	Significant logical errors
1	Fundamental failure in reasoning

Error Types

Type	Description
<code>none</code>	No error (correct answer)
<code>wrong_logic</code>	Correct reasoning structure but logical/computational error
<code>hallucination</code>	Model fabricated information
<code>format_error</code>	Answer format could not be parsed
<code>no_answer</code>	Model did not provide a clear answer
<code>calculation_error</code>	Arithmetic/computational mistake

Report generated by `comprehensive_llm_judge_analyzer.py`